Custom View Settings

Topic 1 - Single Topic

Question #1 *Topic 1*

You are migrating workloads to the cloud. The goal of the migration is to serve customers worldwide as quickly as possible According to local regulations, certain data is required to be stored in a specific geographic area, and it can be served worldwide. You need to design the architecture and deployment for your workloads.

What should you do?

A. Select a public cloud provider that is only active in the required geographic area

B. Select a private cloud provider that globally replicates data storage for fast data access

C. Select a public cloud provider that guarantees data location in the required geographic area

D. Select a private cloud provider that is only active in the required geographic area

Correct Answer: *C*

leyunjohn 4 days, 6 hours ago
The correct is C
upvoted 1 times

Syd 1 week ago
Answer C is correct as it makes sure the data is in the region.
upvoted 2 times

ahsangh 1 week ago
Selected Answer: C
C. Select a public cloud provider that guarantees data location in the required geographic area
upvoted 3 times

kitubha 1 week, 4 days ago
The correct is D
upvoted 1 times

mabkhan 1 week, 3 days ago
why D is correct answer ?
upvoted 1 times

Question #2 *Topic 1*

Your organization needs a large amount of extra computing power within the next two weeks.
After those two weeks, the need for the additional resources will end.
Which is the most cost-effective approach?

A. Use a committed use discount to reserve a very powerful virtual machine

B. Purchase one very powerful physical computer

C. Start a very powerful virtual machine without using a committed use discount

D. Purchase multiple physical computers and scale workload across them

Correct Answer: *C*

leyunjohn 4 days, 5 hours ago
C is correct
upvoted 1 times

Syd 1 week ago
C is correct because we don't know the compute use for 2 weeks and discount is available on 1-3 years not for 2 weeks.

https://cloud.google.com/compute/docs/instances/signing-up-committed-use-discounts
upvoted 4 times

ahsangh 1 week ago
C. Start a very powerful virtual machine without using a committed use discount
upvoted 1 times

kitubha 1 week, 4 days ago
A is correct
upvoted 1 times

Question #3 *Topic 1*

Your organization needs to plan its cloud infrastructure expenditures.
Which should your organization do?

A. Review cloud resource costs frequently, because costs change often based on use

B. Review cloud resource costs annually as part of planning your organizationλ€™s overall budget

C. If your organization uses only cloud resources, infrastructure costs are no longer part of your overall budget

D. Involve fewer people in cloud resource planning than your organization did for on-premises resource planning

Correct Answer: *A*

leyunjohn 4 days, 4 hours ago

Selected Answer: C

Correct Answer: C

upvoted 1 times

leyunjohn 4 days, 5 hours ago

Correct Answer: C

upvoted 1 times

SURJK 6 days, 18 hours ago

Answer A, for C; question doe snot state anything specific whether the usage is Cloud resources only or not

upvoted 4 times

ahsangh 1 week ago

A. Review cloud resource costs frequently, because costs change often based on use

upvoted 2 times

kitubha 1 week, 4 days ago

C is correct

upvoted 1 times

Question #4 *Topic 1*

The operating systems of some of your organizationλ€™s virtual machines may have a security vulnerability.
How can your organization most effectively identify all virtual machines that do not have the latest security update?

A. View the Security Command Center to identify virtual machines running vulnerable disk images

B. View the Compliance Reports Manager to identify and download a recent PCI audit

C. View the Security Command Center to identify virtual machines started more than 2 weeks ago

D. View the Compliance Reports Manager to identify and download a recent SOC 1 audit

Correct Answer: *A*

Maverickps 1 day, 3 hours ago

Selected Answer: A

A sounds Correct

upvoted 1 times

kitubha 1 week, 4 days ago

A is correct

upvoted 2 times

Question #5 *Topic 1*

You are currently managing workloads running on Windows Server for which your company owns the licenses. Your workloads are only needed during working hours, which allows you to shut down the instances during the weekend. Your Windows Server licenses are up for renewal in a month, and you want to optimize your license cost.
What should you do?

A. Renew your licenses for an additional period of 3 years. Renew your licenses for an additional period of 3 years. Negotiate a cost reduction with your current hosting provider wherein infrastructure cost is reduced when workloads are not in use

B. Renew your licenses for an additional period of 2 years. Negotiate a cost reduction by committing to an automatic renewal of the licenses at the end of the 2 year period

C. Migrate the workloads to Compute Engine with a bring-your-own-license (BYOL) model

D. Migrate the workloads to Compute Engine with a pay-as-you-go (PAYG) model

Correct Answer: *D*

leyunjohn 4 days, 4 hours ago

Selected Answer: D

D is correct

upvoted 1 times

shamg 1 week ago

Should be "D" , key word is Your Windows Server licenses are up for renewal in a month, and you want to optimize your license cost. C option will work for "Sole Tenant" node/Bare Metal type server.

upvoted 1 times

zarich 1 week ago

Selected Answer: C

For its C

upvoted 3 times

Syd 1 week ago

D is correct as weekends the VM can be shutdown and don’t have to renew license again.

upvoted 1 times

kitubha 1 week, 4 days ago

D is correct

upvoted 3 times

fpreli 1 week, 5 days ago

D for me, it says "optimize your license cost" and that license will expire in one month, so I would go with a Pay-As-You-Go model

upvoted 3 times

ahsangh 1 week, 5 days ago

Selected Answer: C

C makes sense.

upvoted 2 times

ahsangh 1 week ago

correction: D makes more sense for optimizing cost.

upvoted 2 times

kapomony 1 week, 6 days ago

C for me

upvoted 1 times

Question #6 *Topic 1*

Your organization runs a distributed application in the Compute Engine virtual machines. Your organization needs redundancy, but it also needs extremely fast communication (less than 10 milliseconds) between the parts of the application in different virtual machines.
Where should your organization locate this virtual machines?

A. In a single zone within a single region

B. In different zones within a single region

C. In multiple regions, using one zone per region

D. In multiple regions, using multiple zones per region

Correct Answer: *D*

anverdor 2 days, 4 hours ago
The latency betweem regions is around 3digits and inter-zones 2 digits. So its B
upvoted 1 times

jasbans 2 days, 5 hours ago
D is correct as per Google: for apps such as real-time games, a few milliseconds of latency can have a greater affect on user experience. Deploy these types of apps in multiple regions close to the users.
upvoted 1 times

prekair0 3 days, 2 hours ago
Selected Answer: B
Multi zone is also redundant within the region and it provides the lowest latency.
upvoted 1 times

fpreli 5 days, 22 hours ago
D also for me and +1 for the poorly worded. It asks for redundancy (multi region) and low latency (multi zone), so I would with D
upvoted 1 times

shamg 1 week ago
This is poorly worded Q , if we have to consider redundancy only D is best option.
upvoted 2 times

kitubha 1 week, 4 days ago
B is correct
upvoted 2 times

Question #7 *Topic 1*

You decide to migrate your on-premises environment to the cloud. You need to determine which resource components still need to be assigned ownership.
Which two functions are owned by a public cloud provider? (Choose two.)

A. Hardware maintenance

B. Infrastructure architecture

C. Infrastructure deployment automation

D. Hardware capacity management

E. Fixing application security issues

Correct Answer: *AB*

Mavericks 1 day, 3 hours ago

Selected Answer: AD

A & D because 'Hardware' is the key word here, Which is always the responsibility of the CSP

upvoted 1 times

leyunjohn 4 days, 4 hours ago

A & B are correct

upvoted 1 times

kapomony 5 days, 4 hours ago

A,D for me (harware capaciy ... is provided by the cloud provider for sure, I can't buy/provide HW on the public cloud)

upvoted 2 times

isinha560035 5 days, 8 hours ago

A D for me - any HW maintenance or elasticity(capacity) should be providers duty.

upvoted 1 times

shamg 1 week ago

A,D for me- Hardware maintenance and capacity planning is CSP responsibility going by Shared Responsibility model.

upvoted 3 times

Syd 1 week ago

A & C are owned by cloud provider.

upvoted 1 times

kitubha 1 week, 4 days ago

C & E are correct

upvoted 1 times

fpreli 1 week, 5 days ago

Not sure about this one, I would say AD since customer still needs to pick the infrastructure architecture while hardware capacity is up to Google.

upvoted 2 times

Question #8 *Topic 1*

You are a program manager within a Software as a Service (SaaS) company that offers rendering software for animation studios. Your team needs the ability to allow scenes to be scheduled at will and to be interrupted at any time to restart later. Any individual scene rendering takes less than 12 hours to complete, and there is no service-level agreement (SLA) for the completion time for all scenes. Results will be stored in a global Cloud Storage bucket. The compute resources are not bound to any single geographical location. This software needs to run on Google Cloud in a cost-optimized way.
What should you do?

A. Deploy the application on Compute Engine using preemptible instances

B. Develop the application so it can run in an unmanaged instance group

C. Create a reservation for the minimum number of Compute Engine instances you will use

D. Start more instances with fewer virtual centralized processing units (vCPUs) instead of fewer instances with more vCPUs

Correct Answer: *A*

shamg 1 week ago

A for me -
Can be interrupted at any time to restart later, software needs to run on Google Cloud in a cost-optimized way

upvoted 2 times

Question #9 *Topic 1*

Your manager wants to restrict communication of all virtual machines with internet access; with resources in another network; or with a resource outside Compute
Engine. It is expected that different teams will create new folders and projects in the near future.
How would you restrict all virtual machines from having an external IP address?

A. Define an organization policy at the root organization node to restrict virtual machine instances from having an external IP address

B. Define an organization policy on all existing folders to define a constraint to restrict virtual machine instances from having an external IP address

C. Define an organization policy on all existing projects to restrict virtual machine instances from having an external IP address

D. Communicate with the different teams and agree that each time a virtual machine is created, it must be configured without an external IP address

Correct Answer: *A*
Reference:
https://cloud.google.com/resource-manager/docs/organization-policy/overview

mb0812 2 days, 20 hours ago
Should be A
upvoted 1 times

Question #10 *Topic 1*

Your multinational organization has servers running mission-critical workloads on its premises around the world. You want to be able to manage these workloads consistently and centrally, and you want to stop managing infrastructure.
What should your organization do?

A. Migrate the workloads to a public cloud

B. Migrate the workloads to a central office building

C. Migrate the workloads to multiple local co-location facilities

D. Migrate the workloads to multiple local private clouds

Correct Answer: *A*

akshay_jadhav 3 days, 8 hours ago
A. is correct
upvoted 1 times

Syd 1 week ago
A is correct. Key is manage centrally.
manage these workloads consistently and centrally, and you want to stop managing infrastructure.
upvoted 4 times

kitubha 1 week, 4 days ago
Correct is D
upvoted 3 times

Question #11 *Topic 1*

Your organization stores highly sensitive data on-premises that cannot be sent over the public internet. The data must be processed both on-premises and in the cloud.
What should your organization do?

A. Configure Identity-Aware Proxy (IAP) in your Google Cloud VPC network

B. Create a Cloud VPN tunnel between Google Cloud and your data center

C. Order a Partner Interconnect connection with your network provider

D. Enable Private Google Access in your Google Cloud VPC network

Correct Answer: *C*
After the service provider provisions your connection, you can start passing traffic between your networks by using the service provider's network.
Reference:
https://cloud.google.com/network-connectivity/docs/interconnect/concepts/partner-overview

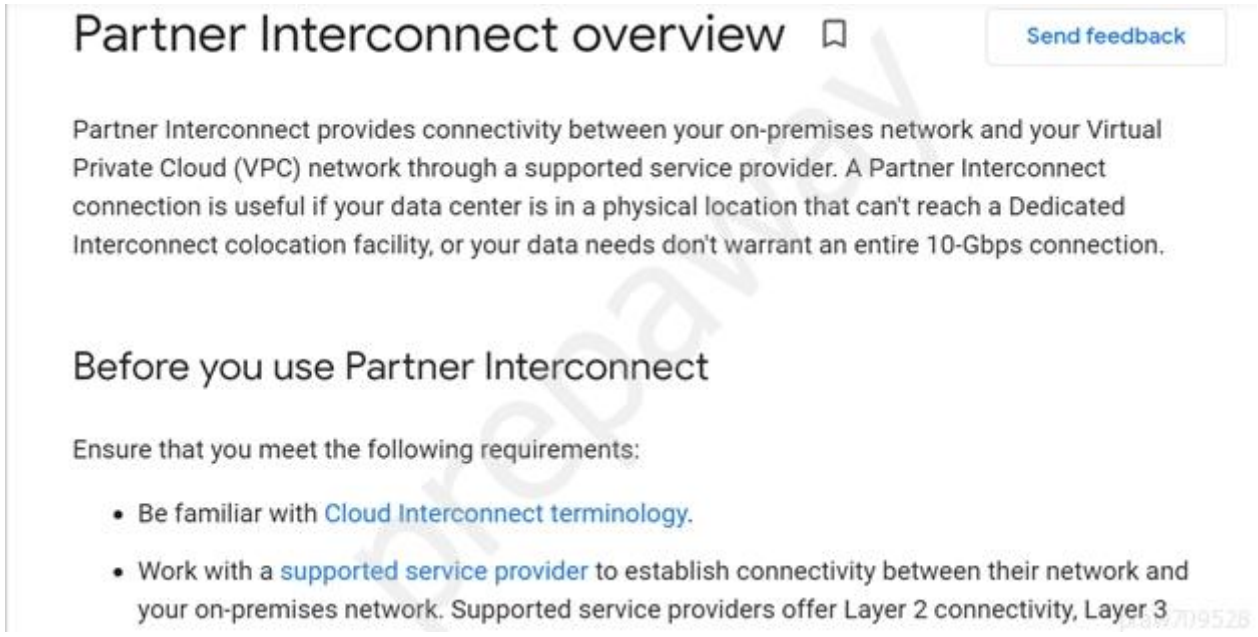
Partner Interconnect overview

Send feedback

Partner Interconnect provides connectivity between your on-premises network and your Virtual Private Cloud (VPC) network through a supported service provider. A Partner Interconnect connection is useful if your data center is in a physical location that can't reach a Dedicated Interconnect colocation facility, or your data needs don't warrant an entire 10-Gbps connection.

Before you use Partner Interconnect

Ensure that you meet the following requirements:

- Be familiar with Cloud Interconnect terminology.
- Work with a supported service provider to establish connectivity between their network and your on-premises network. Supported service providers offer Layer 2 connectivity, Layer 3

akshay_jadhav 3 days, 8 hours ago
option C. because interconnect forms a physical connection between them so data doesnt goes over internet.
upvoted 1 times

shamg 1 week ago
Has to be C
upvoted 2 times

fpreli 1 week, 5 days ago
C for me, Cloud VPN (https://cloud.google.com/network-connectivity/docs/vpn/concepts/overview) goes over the Internet, not acceptable for highly-sensitive data
upvoted 2 times

kiranm2021 1 week, 5 days ago
I also thought B is correct however the answer C given in the examtopics seems to be correct based on this video https://www.youtube.com/watch?v=cKaryf7qp9w
upvoted 1 times

ahsangh 1 week, 5 days ago
Selected Answer: B
B. tunnel
upvoted 3 times

ahsangh 1 week ago
correction: C. interconnect
upvoted 1 times

kapomony 1 week, 6 days ago
B for me
upvoted 3 times

Question #12 *Topic 1*

Your companyλ€™s development team is building an application that will be deployed on Cloud Run. You are designing a CI/CD pipeline so that any new version of the application can be deployed in the fewest number of steps possible using the CI/CD pipeline you are designing. You need to select a storage location for the images of the application after the CI part of your pipeline has built them.
What should you do?

A. Create a Compute Engine image containing the application

B. Store the images in Container Registry

C. Store the images in Cloud Storage

D. Create a Compute Engine disk containing the application

Correct Answer: *B*
Reference:
https://cloud.google.com/container-registry/docs/pushing-and-pulling

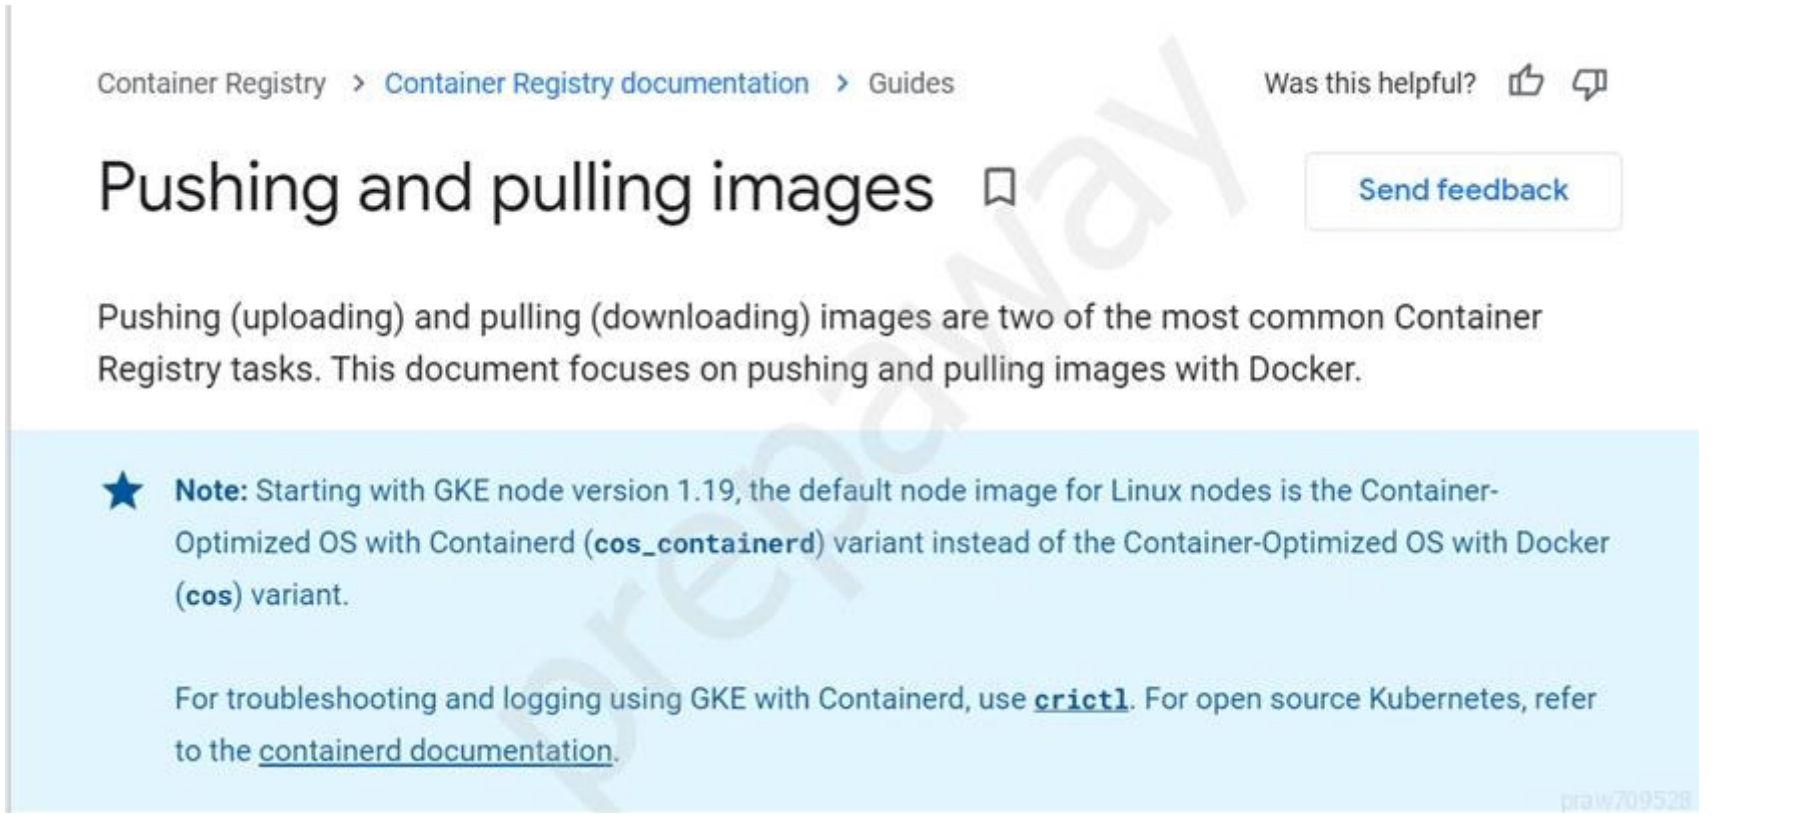
Container Registry > Container Registry documentation > Guides

Was this helpful?

Pushing and pulling images

Send feedback

Pushing (uploading) and pulling (downloading) images are two of the most common Container Registry tasks. This document focuses on pushing and pulling images with Docker.

Note: Starting with GKE node version 1.19, the default node image for Linux nodes is the Container-Optimized OS with Containerd (`cos_containerd`) variant instead of the Container-Optimized OS with Docker (`cos`) variant.

For troubleshooting and logging using GKE with Containerd, use `crictl`. For open source Kubernetes, refer to the containerd documentation.

Question #13 *Topic 1*

Each of the three cloud service models - infrastructure as a service (IaaS), platform as a service (PaaS), and software as a service (SaaS) - offers benefits between flexibility and levels of management by the cloud provider and the customer.
Why would SaaS be the right choice of service model?

A. You want a balance between flexibility for the customer and the level of management by the cloud provider

B. You want to minimize the level of management by the customer

C. You want to maximize flexibility for the customer.

D. You want to be able to shift your emphasis between flexibility and management by the cloud provider as business needs change

Correct Answer: *B*
Benefits of SaaS -
The main benefit of SaaS is that it offloads all infrastructure and application management to the SaaS vendor.
Reference:
https://www.ibm.com/cloud/learn/iaas-paas-saas

Question #14 *Topic 1*

As your organization increases its release velocity, the VM-based application upgrades take a long time to perform rolling updates due to OS boot times. You need to make the application deployments faster.
What should your organization do?

A. Migrate your VMs to the cloud, and add more resources to them

B. Convert your applications into containers

C. Increase the resources of your VMs

D. Automate your upgrade rollouts

Correct Answer: *D*
Reference:
https://cloud.google.com/compute/docs/instance-groups/rolling-out-updates-to-managed-instance-groups


Automatically rolling out updates to instances in a MIG

Send feedback

This document describes how to automatically apply configuration updates to the virtual machine (VM) instances in a managed instance group (MIG).

Compute Engine maintains the VMs in a MIG based on the configuration that you specify in an instance template and optional stateful configuration. From time to time, you might want to update this configuration.

When you set up an automated update, the MIG rolls out a new version of an instance template automatically to all or to a subset of the group's VMs. If you have stateful configuration, the MIG also applies any unapplied per-instance configurations to the corresponding VMs.

You can control the speed of deployment, the level of disruption to your service, and, by using a canary update, the number of instances that the MIG updates with the new template. After you

isinha560035 5 days, 8 hours ago
B for me - I cant justify D - problem to solve is "upgrades take a long time to perform rolling updates due to OS boot times". Thats what we get if we use a COS image.
upvoted 1 times

Terzlightyear 5 days, 21 hours ago
The same question was on Coursera, the right one is the D
upvoted 1 times

fpreli 1 week, 5 days ago
B for me, It says the problem is the OS boot, that is exactly one the key benefits of containerisation, removing the OS boot time.
upvoted 3 times

Question #15 *Topic 1*

Your organization uses Active Directory to authenticate users. Usersλ€™ Google account access must be removed when their Active Directory account is terminated.
How should your organization meet this requirement?

A. Configure two-factor authentication in the Google domain

B. Remove the Google account from all IAM policies

C. Configure BeyondCorp and Identity-Aware Proxy in the Google domain

D. Configure single sign-on in the Google domain

Correct Answer: *D*

JCE 3 days, 10 hours ago
D seems correct
https://cloud.google.com/architecture/identity/federating-gcp-with-active-directory-configuring-single-sign-on
upvoted 1 times

Question #16 *Topic 1*

Your company has recently acquired three growing startups in three different countries. You want to reduce overhead in infrastructure management and keep your costs low without sacrificing security and quality of service to your customers.
How should you meet these requirements?

A. Host all your subsidiariesλ€™ services on-premises together with your existing services.

B. Host all your subsidiariesλ€™ services together with your existing services on the public cloud.

C. Build a homogenous infrastructure at each subsidiary, and invest in training their engineers.

D. Build a homogenous infrastructure at each subsidiary, and invest in hiring more engineers.

Correct Answer: *B*

Question #17 *Topic 1*

What is the difference between Standard and Coldline storage?

A. Coldline storage is for data for which a slow transfer rate is acceptable.

B. Standard and Coldline storage have different durability guarantees.

C. Standard and Coldline storage use different APIs.

D. Coldline storage is for infrequently accessed data.

Correct Answer: *D*

Reference:

https://www.msp360.com/resources/blog/google-cloud-nearline-storage-vs-coldline-vs-standard/

The main characteristics of Coldline are as follows:

- SLA guarantees 99% data availability (of ten thousand hours, data can be offline for up to 100 hours).
- Monthly fee per GB stored is between $0.004-$.0.014, depending on a region.
- The minimum storing period is 90 days. If you delete data earlier, you have to pay for the remaining time. For example, if you had uploaded 100GB of data and then deleted it after 30 days, you need to pay extra $1.4 as an early deletion fee ($0.004 (or $0.014) * 100 GB * 2 months).
- Data retrieval fee is mandatory and costs $0.05 per GB.

Its low price and a long minimum storing period make Nearline the best solution for data that is unlikely to be accessed more than once a year, if ever:

- Data archive.
- Disaster recovery storage.
- Outdated backups storage.

akshay_jadhav 3 days, 7 hours ago

D is correct

upvoted 1 times

Question #18 *Topic 1*

What would provide near-unlimited availability of computing resources without requiring your organization to procure and provision new equipment?

A. Public cloud

B. Containers

C. Private cloud

D. Microservices

Correct Answer: *A*
Reference:
https://cloud.google.com/docs/overview

akshay_jadhav 3 days, 7 hours ago
A. Public cloud
upvoted 1 times

Question #19 *Topic 1*

You are a program manager for a team of developers who are building an event-driven application to allow users to follow one anotherλ€™s activities in the app. Each time a user adds himself as a follower of another user, a write occurs in the real-time database.
The developers will develop a lightweight piece of code that can respond to database writes and generate a notification to let the appropriate users know that they have gained new followers. The code should integrate with other cloud services such as Pub/Sub, Firebase, and Cloud APIs to streamline the orchestration process. The application requires a platform that automatically manages underlying infrastructure and scales to zero when there is no activity.
Which primary compute resource should your developers select, given these requirements?

A. Google Kubernetes Engine

B. Cloud Functions

C. App Engine flexible environment

D. Compute Engine

Correct Answer: *B*

Reference:

https://firebase.google.com/docs/functions/use-cases

Cloud Functions gives developers access to Firebase and Google Cloud events, along with scalable computing power to run code in response to those events. While it's expected that Firebase apps will use Cloud Functions in unique ways to meet their unique requirements, typical use cases might fall into these areas:

- Notify users when something interesting happens.
- Perform database sanitization and maintenance.
- Execute intensive tasks in the cloud instead of in your app.
- Integrate with third-party services and APIs.

Review the use cases and examples for each category that interests you, and then proceed to our Get Started tutorial or to specific how-to guides for authentication events, analytics events, and more. See the `eventType` API reference for the complete list of supported event types.

akshay_jadhav 3 days, 7 hours ago

B. For event based data cloud functions are used

upvoted 1 times

Question #20 *Topic 1*

Your organization is developing an application that will capture a large amount of data from millions of different sensor devices spread all around the world. Your organization needs a database that is suitable for worldwide, high-speed data storage of a large amount of unstructured data. Which Google Cloud product should your organization choose?

A. Firestore

B. Cloud Data Fusion

C. Cloud SQL

D. Cloud Bigtable

Correct Answer: *D*

Reference:
https://cloud.google.com/bigtable

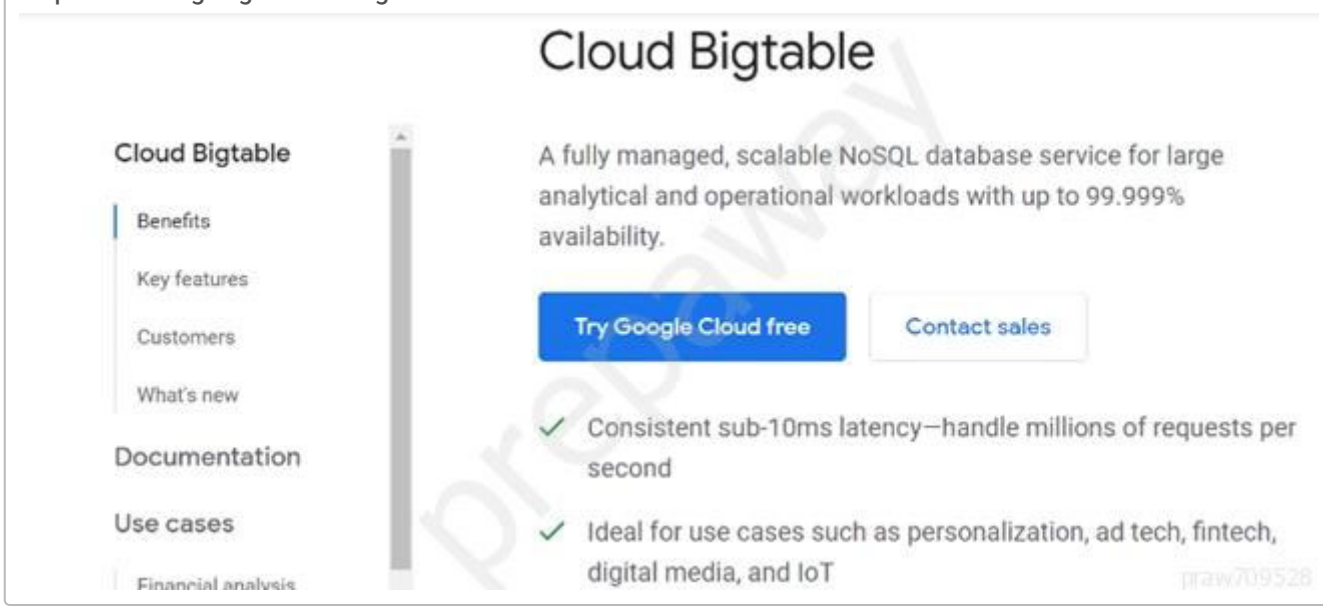

akshay_jadhav 3 days, 6 hours ago
d is correct ans. for iot based unstructured data we use bigtable.
upvoted 1 times

Syd 4 days, 20 hours ago
C correct Answer. Key is devices=IOT

https://cloud.google.com/bigtable
upvoted 1 times

Nidie 19 hours, 33 minutes ago
It is unstructured data. It cannot be C
upvoted 1 times

akshay_jadhav 3 days, 6 hours ago
your explanation is correct but its option D
upvoted 1 times

AshSM 5 days, 14 hours ago
Option is A as Firestore is globally distributed. Option D, Bigtable, is for Hadoop and locally accessed workloads that need high performance.
upvoted 2 times

Question #21 *Topic 1*

Your organization needs to build streaming data pipelines. You donλ€™t want to manage the individual servers that do the data processing in the pipelines. Instead, you want a managed service that will automatically scale with the amount of data to be processed.
Which Google Cloud product or feature should your organization choose?

A. Pub/Sub

B. Dataflow

C. Data Catalog

D. Dataprep by Trifacta

Correct Answer: *B*
Reference:
https://cloud.google.com/dataflow/docs/guides/deploying-a-pipeline

Deploying a pipeline

Send feedback

★ This document explains in detail how Dataflow deploys and runs a pipeline, and covers advanced topics like optimization and load balancing. If you are looking for a step-by-step guide on how to create and deploy your first pipeline, use Dataflow's quickstarts for Java, Python or templates.

After you construct and test your Apache Beam pipeline, you can use the Dataflow managed service to deploy and execute it. Once on the Dataflow service, your pipeline code becomes a Dataflow job.

The Dataflow service fully manages Google Cloud services such as Compute Engine and Cloud Storage to run your Dataflow job, automatically spinning up and tearing down the necessary resources. The Dataflow service provides visibility into your job through tools like the Dataflow monitoring interface and the Dataflow command-line interface.

Question #22 *Topic 1*

Your organization is building an application running in Google Cloud. Currently, software builds, tests, and regular deployments are done manually, but you want to reduce work for the team. Your organization wants to use Google Cloud managed solutions to automate your build, testing, and deployment process.
Which Google Cloud product or feature should your organization use?

A. Cloud Scheduler

B. Cloud Code

C. Cloud Build

D. Cloud Deployment Manager

Correct Answer: *C*
Deploy your application to App Engine using the gcloud app deploy command. This command automatically builds a container image by using the Cloud
Build service and then deploys that image to the App Engine flexible environment.
Reference:
https://cloud.google.com/appengine/docs/flexible/nodejs/testing-and-deploying-your-app

Question #23 *Topic 1*

Which Google Cloud product can report on and maintain compliance on your entire Google Cloud organization to cover multiple projects?

A. Cloud Logging

B. Identity and Access Management

C. Google Cloud Armor

D. Security Command Center

Correct Answer: *B*

kshitij1992 4 days, 12 hours ago

Selected Answer: D

Answer is D

upvoted 1 times

Syd 4 days, 20 hours ago

Answer D is correct.

https://cloud.google.com/security-command-center

upvoted 1 times

kapomony 5 days, 3 hours ago

is SCC

upvoted 1 times

SURJK 6 days, 10 hours ago

Answer is D: SCC - Gain centralized visibility and control
Discover misconfigurations and vulnerabilities
Report on and maintain compliance
Detect threats targeting your Google Cloud assets

upvoted 4 times

ahsangh 1 week ago

Selected Answer: B

B. Identity and Access Management

https://cloud.google.com/docs/enterprise/best-practices-for-enterprise-organizations

upvoted 1 times

isinha560035 1 week, 5 days ago

Its SCC, can you correct suggested answer please.

upvoted 2 times

kiranm2021 1 week, 5 days ago

Is this correct ? Isn't this Compliance Manager . Not sure how IAM manages compliance reports.

upvoted 1 times

Question #24 *Topic 1*

Your organization needs to establish private network connectivity between its on-premises network and its workloads running in Google Cloud.
You need to be able to set up the connection as soon as possible.
Which Google Cloud product or feature should you use?

A. Cloud Interconnect

B. Direct Peering

C. Cloud VPN

D. Cloud CDN

Correct Answer: *A*

A_A_AB 1 day, 11 hours ago

Selected Answer: C

C makes more sence as creating VPN Connection is almost instant, thus faster than any other option.

upvoted 1 times

anverdor 2 days, 4 hours ago

It's VPN because Interconnect needs more time to be ready

upvoted 1 times

kshitij1992 4 days, 12 hours ago

Selected Answer: A

Answer is A
https://cloud.google.com/network-connectivity/docs/interconnect/concepts/overview

upvoted 1 times

SURJK 3 days, 14 hours ago

Isn't setting Cloud VPN faster than Cloud Interconnect ?

upvoted 1 times

SURJK 6 days, 10 hours ago

Answer is C: Cloud VPN

upvoted 3 times

Question #25 *Topic 1*

Your organization is developing a mobile app and wants to select a fully featured cloud-based compute platform for it.
Which Google Cloud product or feature should your organization use?

A. Google Kubernetes Engine

B. Firebase

C. Cloud Functions

D. App Engine

Correct Answer: *B*

Syd 4 days, 20 hours ago
B is correct. Key is mobile app

https://cloud.google.com/firestore/docs/client/get-firebase
upvoted 1 times

SURJK 6 days, 9 hours ago
Answer is B considering GCP bringing in Firebase as Mobile platform. but the question asked is about Compute platform, Is Firebase Compute platform ?
upvoted 1 times

fpreli 5 days, 22 hours ago
It seems yes, "Firebase is Google's mobile development platform" from their documentation
upvoted 1 times

cyqgz_36 6 days, 17 hours ago
B is correct
upvoted 1 times

kapomony 1 week ago
"mobile" is the key word, so "firestore" is correct
upvoted 1 times

ahsangh 1 week, 5 days ago
Selected Answer: D
D. App Engine
upvoted 4 times

mabkhan 1 week, 2 days ago
why Answer D is correct?
upvoted 1 times

ahsangh 1 week ago
https://cloud.google.com/appengine
upvoted 1 times

A_A_AB 1 day, 11 hours ago
https://cloud.google.com/firebase
upvoted 1 times

Question #26 *Topic 1*

Your company has been using a shared facility for data storage and will be migrating to Google Cloud. One of the internal applications uses Linux custom images that need to be migrated.
Which Google Cloud product should you use to maintain the custom images?

A. App Engine flexible environment

B. Compute Engine

C. App Engine standard environment

D. Google Kubernetes Engine

Correct Answer: *B*
Reference:
https://cloud.google.com/compute/docs/images/create-delete-deprecate-private-images

SURJK 6 days, 9 hours ago
Compute Engine
upvoted 1 times

ahsangh 1 week, 5 days ago
Can we import custom images into GKE?
upvoted 1 times

Question #27 *Topic 1*

Your organization wants to migrate its data management solutions to Google Cloud because it needs to dynamically scale up or down and to run transactional
SQL queries against historical data at scale. Which Google Cloud product or service should your organization use?

A. BigQuery

B. Cloud Bigtable

C. Pub/Sub

D. Cloud Spanner

Correct Answer: *D*
Reference:
https://cloud.google.com/terms/services

Cloud Spanner: Cloud Spanner is a fully-managed, mission-critical relational database service. It is designed to provide a scalable online transaction processing (OLTP) database with high availability and strong consistency at global scale.

SURJK 6 days, 9 hours ago
Answer: D --> Q? states transactional SQL, so its Spanner
upvoted 2 times

Question #28 *Topic 1*

Your organization needs to categorize objects in a large group of static images using machine learning. Which Google Cloud product or service should your organization use?

A. BigQuery ML

B. AutoML Video Intelligence

C. Cloud Vision API

D. AutoML Tables

Correct Answer: *D*
Reference:
https://towardsdatascience.com/building-image-detection-with-google-cloud-automl-8b9cf2b8074b?gi=298eab52b978

isinha560035 5 days, 6 hours ago
C for me - its AutoML vision - https://cloud.google.com/vision/automl/docs/create-datasets
upvoted 2 times

AshSM 5 days, 13 hours ago
If it was standard image detection (like people, automobiles) then Vision API makes sense. The question is asking for object detection using machine learning, so AutoML makes more sense.
upvoted 1 times

mvdm11111 1 week, 5 days ago
this should be C Vision API
upvoted 4 times

Question #29 *Topic 1*

Your organization runs all its workloads on Compute Engine virtual machine instances. Your organization has a security requirement: the virtual machines are not allowed to access the public internet. The workloads running on those virtual machines need to access BigQuery and Cloud Storage, using their publicly accessible interfaces, without violating the security requirement.
Which Google Cloud product or feature should your organization use?

A. Identity-Aware Proxy

B. Cloud NAT (network address translation)

C. VPC internal load balancers

D. Private Google Access

Correct Answer: *B*
Reference:
https://cloud.google.com/nat

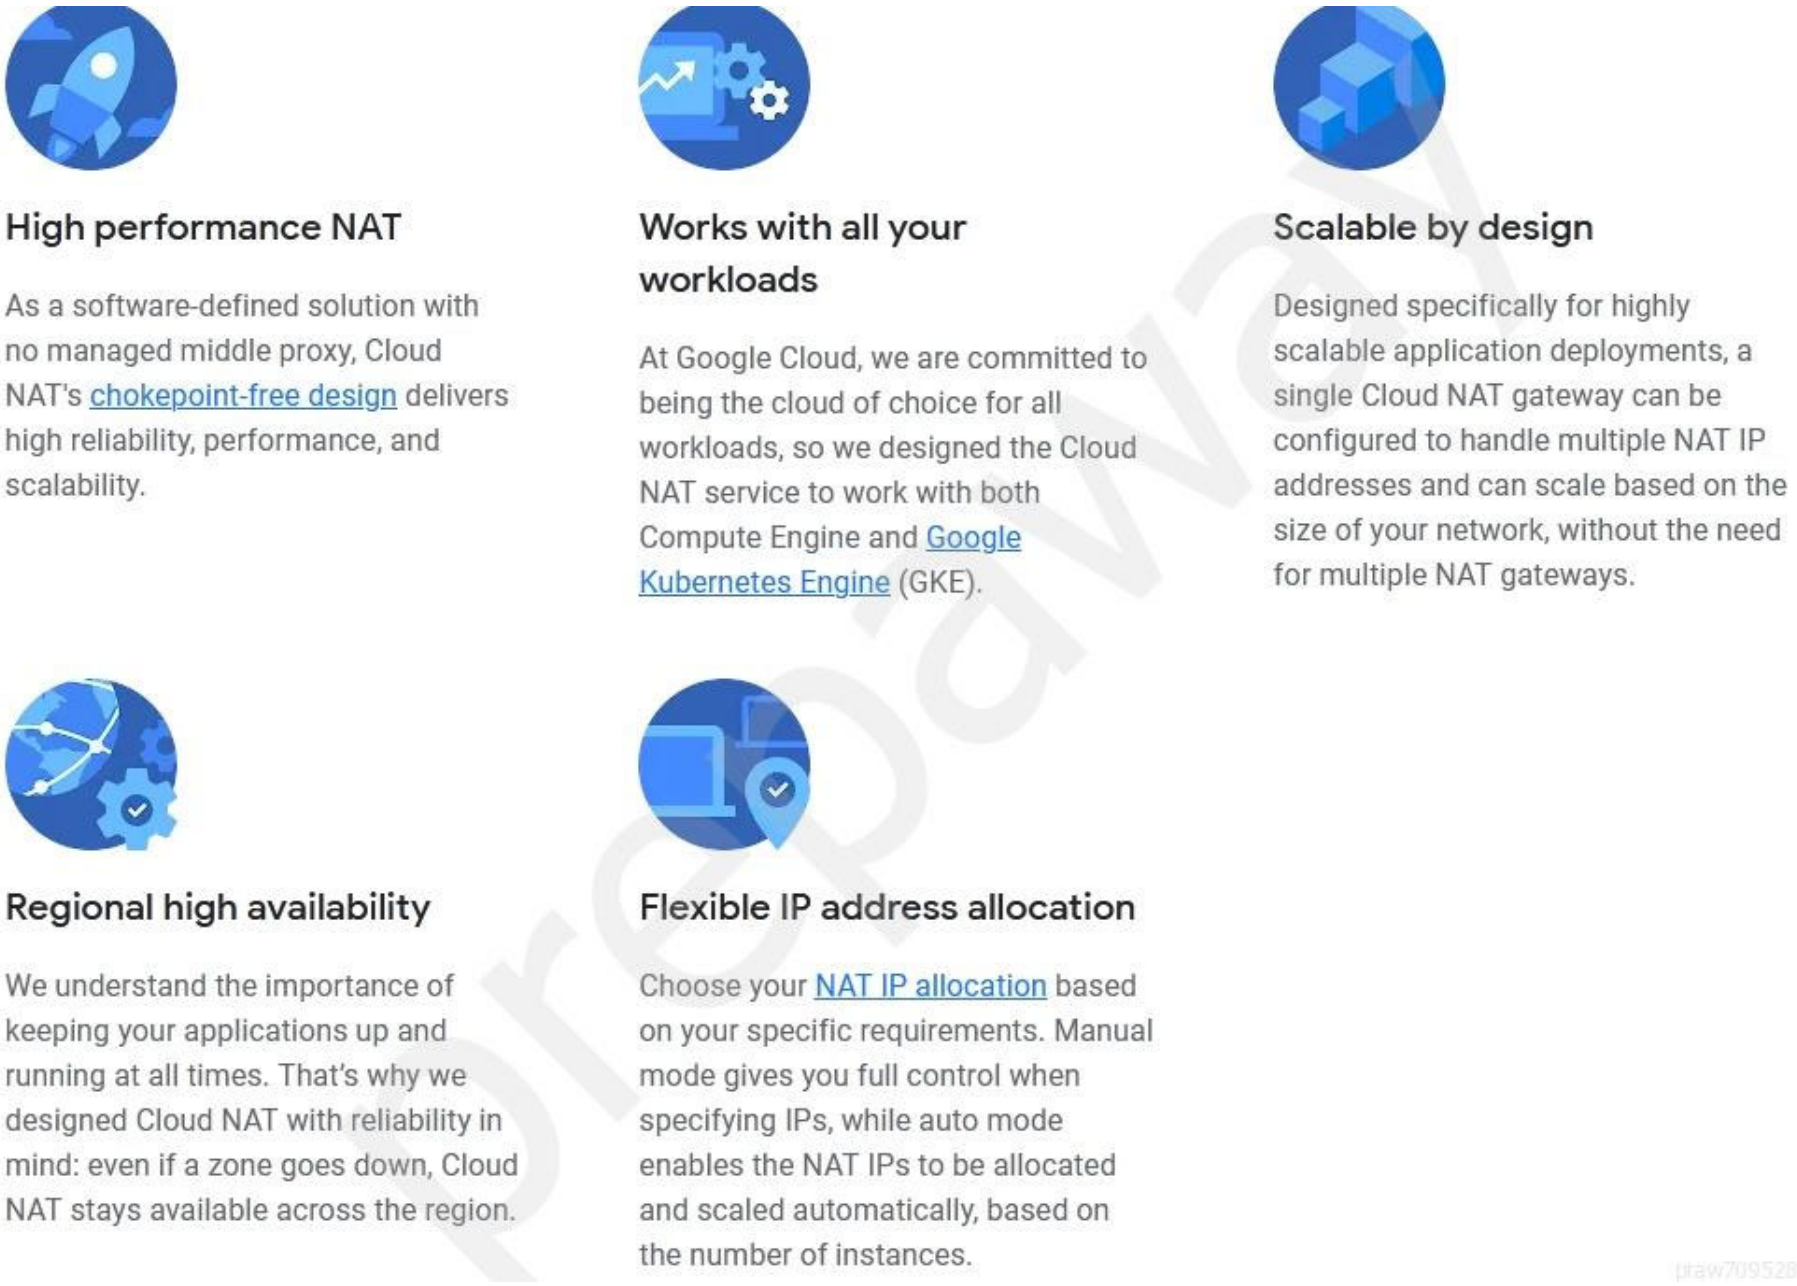

A_A_AB 1 day, 10 hours ago

Selected Answer: D
Agree with fpreli, the answer is D.
According to the Google Documents >>> By default, when a Compute Engine VM lacks an external IP address assigned to its network interface, it can only send packets to other internal IP address destinations. You can allow these VMs to connect to the set of external IP addresses used by Google APIs and services by enabling Private Google Access on the subnet used by the VM's network interface.
Besides, the security rules say no Internet access while NAt provite internet access.
Google Private Access is like AWS VPC endpoint where you access GCP Public services without using Public Internet.

upvoted 1 times

SURJK 6 days, 9 hours ago

Answer B:
Google Cloud's managed Network Address Translation service, enables you to provision your application instances without public IP addresses while also allowing them to access the internet

upvoted 1 times

fpreli 1 week, 5 days ago

I think this one is D since (from GCP doc) "connect to the set of external IP addresses used by Google APIs and services by enabling Private Google Access". On the other hand, Cloud NAT automatically enable Private Google Access, not sure.

upvoted 2 times

Question #30 *Topic 1*

Which Google Cloud product is designed to reduce the risks of handling personally identifiable information (PII)?

A. Cloud Storage

B. Google Cloud Armor

C. Cloud Data Loss Prevention

D. Secret Manager

Correct Answer: *C*

Reference:

https://cloud.google.com/blog/products/gcp/take-charge-of-your-sensitive-data-with-the-cloud-dlp-api

Question #31 *Topic 1*

Your organization is migrating to Google Cloud. As part of that effort, it needs to move terabytes of data from on-premises file servers to Cloud Storage. Your organization wants the migration process to be automated and to be managed by Google. Your organization has an existing Dedicated Interconnect connection that it wants to use. Which Google Cloud product or feature should your organization use?

A. Storage Transfer Service

B. Migrate for Anthos

C. BigQuery Data Transfer Service

D. Transfer Appliance

Correct Answer: *A*

Reference:

https://cloud.google.com/architecture/migration-to-google-cloud-transferring-your-large-datasets

Where you're moving data from	Scenario	Suggested products
Another cloud provider (for example, Amazon Web Services or Microsoft Azure) to Google Cloud	—	Storage Transfer Service
Cloud Storage to Cloud Storage (two different buckets)	—	Storage Transfer Service
Your private data center to Google Cloud	Enough bandwidth to meet your project deadline for less than 1 TB of data	`gsutil`
Your private data center to Google Cloud	Enough bandwidth to meet your project deadline for more than 1 TB of data	Storage Transfer Service for on-premises data
Your private data center to Google Cloud	Not enough bandwidth to meet your project deadline	Transfer Appliance

Question #32 *Topic 1*

Your organization needs to analyze data in order to gather insights into its daily operations. You only want to pay for the data you store and the queries you perform. Which Google Cloud product should your organization choose for its data analytics warehouse?

A. Cloud SQL

B. Dataproc

C. Cloud Spanner

D. BigQuery

Correct Answer: *A*

Syd 4 days, 19 hours ago
D correct answer. Key pay per query used.

https://cloud.google.com/bigquery/pricing
upvoted 1 times

fpreli 1 week, 5 days ago
D for me as well
upvoted 3 times

kapomony 1 week, 5 days ago
BigQuey for me
upvoted 2 times

ahsangh 1 week, 5 days ago
Selected Answer: D
D. BigQuery > SQL
upvoted 1 times

sarangdh 1 week, 6 days ago
Selected Answer: D
for all database analytics and datawarehouse , its should be Bigquery , not Cloud SQL
upvoted 3 times

Question #33 *Topic 1*

Your organization wants to run a container-based application on Google Cloud. This application is expected to increase in complexity. You have a security need for fine-grained control of traffic between the containers. You also have an operational need to exercise fine-grained control over the applicationλ€™s scaling policies.
What Google Cloud product or feature should your organization use?

A. Google Kubernetes Engine cluster

B. App Engine

C. Cloud Run

D. Compute Engine virtual machines

Correct Answer: *C*
Reference:
https://cloud.google.com/blog/products/serverless/cloud-run-gets-enterprise-grade-network-security-with-vpc-sc

kimkimkimkimkim 3 days, 5 hours ago

Selected Answer: A

I think the answer should be A because it says that there is a need for "fine-grained control".
Check out this brief comparison:
https://cloud.google.com/blog/products/containers-kubernetes/when-to-use-google-kubernetes-engine-vs-cloud-run-for-containers

upvoted 3 times

Syd 4 days, 19 hours ago

C correct answer. Security for applications and operation policies.
https://cloud.google.com/run#all-features

upvoted 1 times

fpreli 5 days, 22 hours ago

I would say A. GKE seems a better fit since the requirement is for "security need for fine-grained control of traffic between the containers" and "fine-grained control over scaling policies". Such level of control is easier on GKE that Cloud Run.

upvoted 4 times

Question #34 *Topic 1*

Which Google Cloud product or feature makes specific recommendations based on security risks and compliance violations?

A. Google Cloud firewalls

B. Security Command Center

C. Cloud Deployment Manager

D. Google Cloud Armor

Correct Answer: *B*

Reference:

https://cloud.google.com/security-command-center

Asset discovery and inventory

Discover and view your assets in near-real time across App Engine, BigQuery, Cloud SQL, Cloud Storage, Compute Engine, Cloud Identity and Access Management, Google Kubernetes Engine, and more. Review historical discovery scans to identify new, modified, or deleted assets.

Threat prevention

Understand the security state of your Google Cloud assets. Uncover common web application vulnerabilities such as cross-site scripting or outdated libraries in your web applications running on App Engine, GKE, and Compute Engine. Quickly resolve misconfigurations by clicking directly on the impacted resource and following the prescribed steps on how to fix it.

Threat detection

Detect threats using logs running in Google Cloud at scale. Detect some of the most common container attacks, including suspicious binary, suspicious library, and reverse shell.

SURJK 6 days, 1 hour ago

Answer: B

upvoted 1 times

Question #35 *Topic 1*

Which Google Cloud product gives you a consistent platform for multi-cloud application deployments and extends other Google Cloud services to your environment?

A. Google Kubernetes Engine

B. Virtual Public Cloud

C. Compute Engine

D. Anthos

Correct Answer: *D*

kshitij1992 4 days, 12 hours ago

Selected Answer: D

Answer is D

upvoted 1 times

SURJK 6 days, 1 hour ago

Answer: D

upvoted 1 times

Question #36 *Topic 1*

Your organization is developing an application that will manage payments and online bank accounts located around the world. The most critical requirement for your database is that each transaction is handled consistently. Your organization anticipates almost unlimited growth in the amount of data stored.
Which Google Cloud product should your organization choose?

A. Cloud SQL

B. Cloud Storage

C. Firestore

D. Cloud Spanner

Correct Answer: *D*
Features of Cloud Spanner -
Reference:
https://k21academy.com/google-cloud/cloud-sql-vs-cloud-spanner/

Google Cloud SQL is a fully managed service offered by Google Cloud Platform. **Google Cloud SQL is a MySQL database inside Google Cloud**. There is no need to install, maintain and create admin accounts because it is fully managed by Google Cloud. It helps you create, modify, configure and utilize a relational database, same as MySQL.

Google sends constant updates and adds new features to its services to fulfill the business requirements of its users.

Let's have a look at the improvements done in Google Cloud SQL.

- Google deliberately increased the storage space to **100 Gigabytes** which was 10 Gigabytes earlier.
- The modified version is loaded with the capacity of **16GB RAM** to run instances hassle-free.
- Increased RAM has helped users to keep **four times more cache** than earlier.
- Now Google provides both **replicated and non-replicated databases**.

A_A_AB 1 day, 10 hours ago

Selected Answer: C
Cloud Spanner key feature >>> It provides strong transactional consistency.
upvoted 1 times

Question #37 *Topic 1*

Your organization wants an economical solution to store data such as files, graphical images, and videos and to access and share them securely. Which Google Cloud product or service should your organization use?

A. Cloud Storage

B. Cloud SQL

C. Cloud Spanner

D. BigQuery

Correct Answer: *A*

Reference:

https://www.myamberlife.com/learn/top-ways-to-store-your-digital-files/

Desktop Storage

Despite many external solutions for digital files, some people still store their photos, videos, and content files on their desktop or laptop. The only problem with this method is that your computer can quickly become cluttered with thousands of files. It slows your prized piece of hardware (computer) down.

When you want to find a digital file you probably *expect* that file to come flying up on your screen in an instant. Yet — anyone who keeps a lot of photos on a computer knows it can take minutes, sometimes hours, to find one – even if you keep it on your desktop. It's just not all that convenient to store things this way. Most importantly, just storing these digital files on a desktop leaves them vulnerable to viruses, damage, or theft. Folks who rely on this also generally don't have a back-up plan.

A_A_AB 1 day, 10 hours ago

Selected Answer: A

If you're familiar with AWS, Google Storage is GCP's version of AWS Simple Storage Service (S3) and an S3 bucket would be equivalent to a Google Storage bucket across the two clouds.

upvoted 1 times

Question #38 *Topic 1*

Your organization wants to predict the behavior of visitors to its public website. To do that, you have decided to build a machine learning model. Your team has database-related skills but only basic machine learning skills, and would like to use those database skills.
Which Google Cloud product or feature should your organization choose?

A. BigQuery ML

B. LookML

C. TensorFlow

D. Cloud SQL

Correct Answer: *A*
Reference:
https://cloud.google.com/architecture/predicting-customer-propensity-to-buy

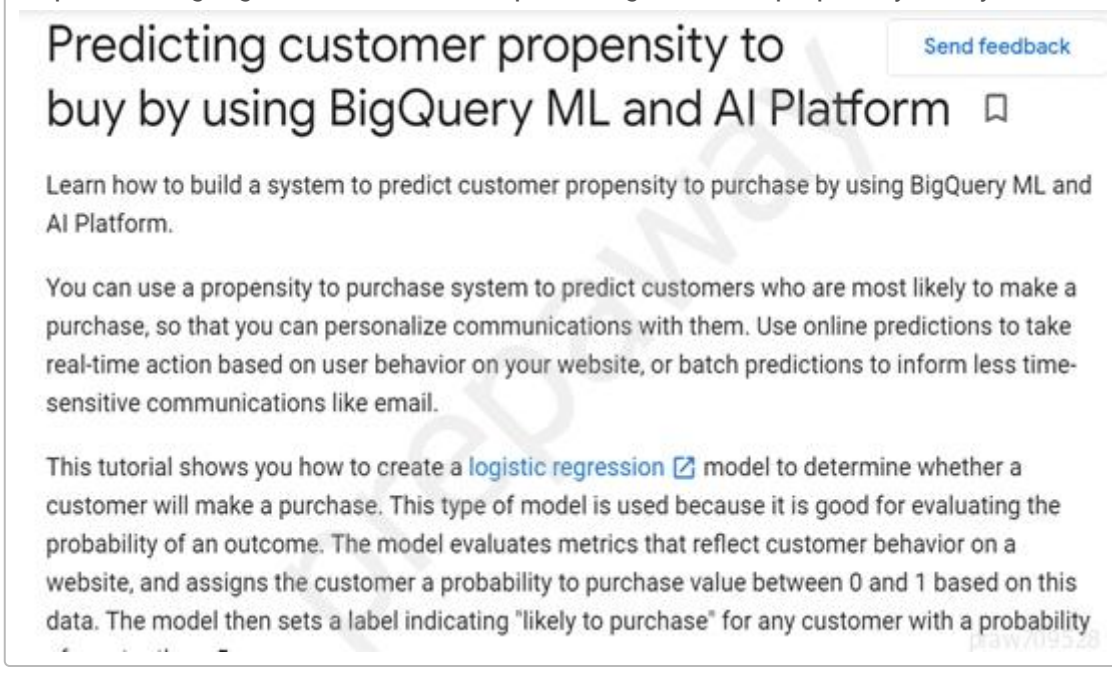

Predicting customer propensity to buy by using BigQuery ML and AI Platform

Send feedback

Learn how to build a system to predict customer propensity to purchase by using BigQuery ML and AI Platform.

You can use a propensity to purchase system to predict customers who are most likely to make a purchase, so that you can personalize communications with them. Use online predictions to take real-time action based on user behavior on your website, or batch predictions to inform less time-sensitive communications like email.

This tutorial shows you how to create a logistic regression model to determine whether a customer will make a purchase. This type of model is used because it is good for evaluating the probability of an outcome. The model evaluates metrics that reflect customer behavior on a website, and assigns the customer a probability to purchase value between 0 and 1 based on this data. The model then sets a label indicating "likely to purchase" for any customer with a probability

Question #39 *Topic 1*

Your organization needs to restrict access to a Cloud Storage bucket. Only employees who are based in Canada should be allowed to view the contents.
What is the most effective and efficient way to satisfy this requirement?

A. Deploy the Cloud Storage bucket to a Google Cloud region in Canada

B. Configure Google Cloud Armor to allow access to the bucket only from IP addresses based in Canada

C. Give each employee who is based in Canada access to the bucket

D. Create a group consisting of all Canada-based employees, and give the group access to the bucket

Correct Answer: *D*

ahsangh 1 week, 4 days ago

Selected Answer: B

B. why not let cloud armor filter by IP addresses?

upvoted 1 times

ThrowAwayBby 1 week ago

I guess it's because you can use your own private VPN to access the Canada-only bucket from anywhere in the world.

upvoted 1 times

Question #40 *Topic 1*

Your organization is moving an application to Google Cloud. As part of that effort, it needs to migrate the applicationλ€™s working database from another cloud provider to Cloud SQL. The database runs on the MySQL engine. The migration must cause minimal disruption to users. Data must be secured while in transit.
Which should your organization use?

A. BigQuery Data Transfer Service

B. MySQL batch insert

C. Database Migration Service

D. Cloud Composer

Correct Answer: *C*

Reference:

https://aws.amazon.com/dms/

Question #41 *Topic 1*

Your organization is developing and deploying an application on Google Cloud. Tracking your Google Cloud spending needs to stay as simple as possible.
What should you do to ensure that workloads in the development environment are fully isolated from production workloads?

A. Apply a unique tag to development resources

B. Associate the development resources with their own network

C. Associate the development resources with their own billing account

D. Put the development resources in their own project

Correct Answer: *D*

Question #42 *Topic 1*

Your company is running the majority of its workloads in a co-located data center. The workloads are running on virtual machines (VMs) on top of a hypervisor and use either Linux or Windows server editions. As part of your companyλ€™s transformation strategy, you need to modernize workloads as much as possible by adopting cloud-native technologies. You need to migrate the workloads into Google Cloud.
What should you do?

A. Export the VMs into VMDK format, and import them into Compute Engine

B. Export the VMs into VMDK format, and import them into Google Cloud VMware Engine

C. Migrate the workloads using Migrate for Compute Engine

D. Migrate the workloads using Migrate for Anthos

Correct Answer: *D*

fpreli 5 days, 21 hours ago
Agree with C, Compute Engine. From GCP documentation "For the workloads that will benefit from containers, Migrate for Anthos and GKE delivers a fast, smooth path to modernization. For other workloads that are better suited as a VM, simply move them as is with Migrate for Compute Engine "
upvoted 1 times

fpreli 5 days, 21 hours ago
But, right, it says "cloud-native" and "modernize", so probably Migrate for Anthos is correct.
upvoted 1 times

kiranm2021 1 week, 5 days ago
may be because it is asking cloud native Anthos is selected ?
upvoted 2 times

kiranm2021 1 week, 5 days ago
How can this be Migrate for Anthos ? Migrate for compute engine is the right answer. Please clarify
upvoted 3 times

Question #43 *Topic 1*

Your organization is running all its workloads in a private cloud on top of a hypervisor. Your organization has decided it wants to move to Google Cloud as quickly as possible. Your organization wants minimal changes to the current environment, while using the maximum amount of managed services Google offers.
What should your organization do?

A. Migrate the workloads to Google Cloud VMware Engine

B. Migrate the workloads to Compute Engine

C. Migrate the workloads to Bare Metal Solution

D. Migrate the workloads to Google Kubernetes Engine

Correct Answer: *B*
Migrate for Compute Engine enables you to lift and shift workloads at scale to Google Cloud Compute Engine with minimal changes and risk.
Reference:
https://dataintegration.info/simplify-vm-migrations-with-migrate-for-compute-engine-as-a-service

> Customers across industries are migrating applications at scale to the cloud. Migrate for Compute Engine enables you to lift and shift workloads at scale to Google Cloud Compute Engine with minimal changes and risk. Recently, we delivered our latest release of Migrate for Compute Engine – Version 5 or V5. This version delivers a Google-managed cloud service that enables simple, frictionless, and large-scale enterprise migrations of virtual machines to Google Compute Engine with minimal downtime and risk. API-driven and integrated into your Google Cloud console for ease of use, this service uses agent-less replication to copy data without manual intervention and without VPN requirements. It also enables you to launch non-disruptive validations of your VMs prior to cutover. Rapidly migrate a single application or execute a sprint with hundred systems using migration groups with confidence.

Question #44 *Topic 1*

Your organization is releasing its first publicly available application in Google Cloud. The application is critical to your business and customers and requires a 2- hour SLA.
How should your organization set up support to minimize costs?

A. Enroll in Premium Support

B. Enroll in Enhanced Support

C. Enroll in Standard Support

D. Enroll in Basic Support

Correct Answer: *B*
Reference:
https://www.secureauth.com/enhanced-support-offering/

SecureAuth is dedicated to providing the **industry-leading** enhanced **support** ensuring the long term success of your SecureAuth SaaS IAM deployment.

While our basic support offers industry leading coverage and response times for some customers, SecureAuth protects critical applications meaning delays and extended downtime is simply not an option. For these customers our **Enhanced Support** offerings provide **24*7 coverage** and the most responsive and **complete SLA**'s available. That's why we offer three different levels of support, so you can choose the level of support that best works for your needs.

Question #45 *Topic 1*

Your organization offers public mobile apps and websites. You want to migrate to a Google Cloud-based solution for checking and maintaining your usersλ€™ usernames and passwords and controlling their access to different resources based on their identity.
Which should your organization choose?

A. VPN tunnels

B. Identity Platform

C. Compute Engine firewall rules

D. Private Google Access

Correct Answer: *B*
An identity platform is a modern solution for managing the identities of users and devices in a centralized fashion.
Reference:
https://www.okta.com/blog/2021/07/what-is-an-identity-platform/#:~:text=An%20identity%20platform%20is%20a,%2C%20integrations%2C%20and%20platform%20services

An identity platform is a modern solution for managing the identities of users and devices in a centralized fashion. It enables organizations to securely authorize workforce and customer users to access their ecosystem using access management tools, programmable components, integrations, and platform services.

Today's organizations have a wide range of identity requirements. As they expand, embrace new innovations, and meet new customer demands, they need an identity solution that can grow with them. And as they grow, these organizations are also looking for offerings that centralize and consolidate identity, reducing the need for multiple access management, governance, and authentication products that don't necessarily talk to each other.

In short, although the concept of identity platforms is still new, the growth of disruptive technologies and highly personalized products and services has made a platform approach the key for effectively unifying identity management. But how did we get here?

Question #46 *Topic 1*

Which Google Cloud service or feature lets you build machine learning models using Standard SQL and data in a data warehouse?

A. BigQuery ML

B. TensorFlow

C. AutoML Tables

D. Cloud Bigtable ML

Correct Answer: *A*

BigQuery ML lets you create and execute machine learning models in BigQuery using standard SQL queries.

Reference:

https://cloud.google.com/bigquery-ml/docs/introduction#:~:text=BigQuery%20ML%20lets%20you%20create,the%20need%20to%20move%20data

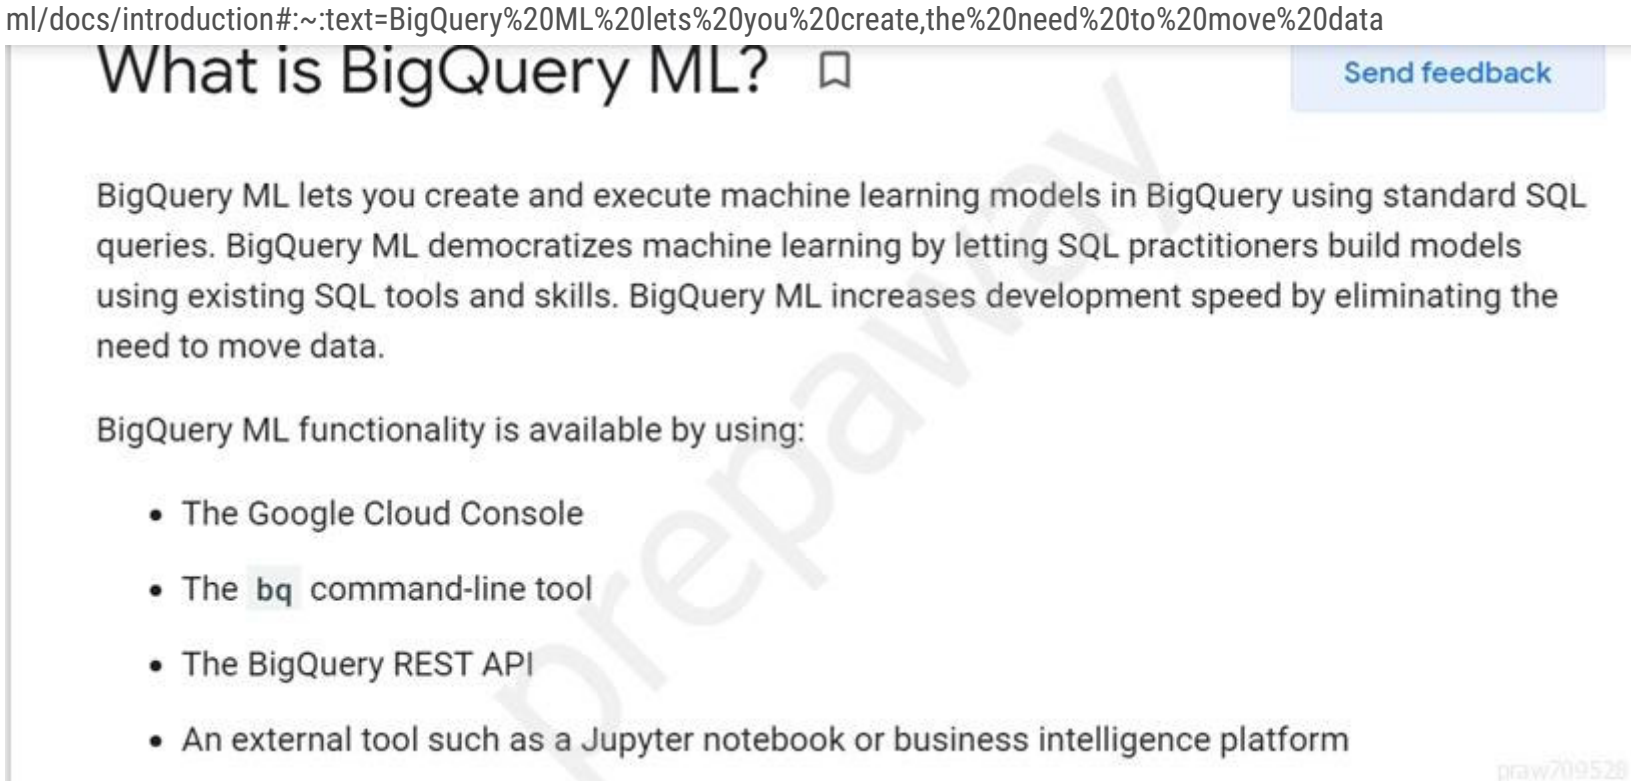

What is BigQuery ML?

Send feedback

BigQuery ML lets you create and execute machine learning models in BigQuery using standard SQL queries. BigQuery ML democratizes machine learning by letting SQL practitioners build models using existing SQL tools and skills. BigQuery ML increases development speed by eliminating the need to move data.

BigQuery ML functionality is available by using:

- The Google Cloud Console
- The `bq` command-line tool
- The BigQuery REST API
- An external tool such as a Jupyter notebook or business intelligence platform

Question #47 *Topic 1*

Your organization runs an application on virtual machines in Google Cloud. This application processes incoming images. This activity takes hours to create a result for each image. The workload for this application normally stays at a certain baseline level, but at regular intervals it spikes to a much greater workload.
Your organization needs to control the cost to run this application.
What should your organization do?

A. Purchase committed use discounts for the baseline load

B. Purchase committed use discounts for the expected spike load

C. Leverage sustained use discounts for your virtual machines

D. Run the workload on preemptible VM instances

Correct Answer: *C*
The idea of the Sustained Use discount is that the longer you run a VM instance in any given month, the bigger discount you will get from the list price.
Reference:
https://www.parkmycloud.com/blog/google-sustained-use-discounts/

price. The following shows the incremental discount, and its c
impact on a hypothetical $100/month VM instance, where the
are against the baseline 730-hour month.

Usage level		Hours usage before price break	discount	Base cost/mo
From	To			$ 100.00
0%	25%	-	0%	$ 25.00
25%	50%	182.5	20%	$ 20.00
50%	75%	365.0	40%	$ 15.00
75%	100%	547.5	60%	$ 10.00
			Net cost	$ 70.00

fpreli 1 week, 4 days ago
A for me, with CUD you get a greater discount and it makes totally sense to buy a CUD for the baseline traffic since it will always be consumed
upvoted 2 times

ahsangh 1 week, 4 days ago
Selected Answer: A
why not use committed use discounts for the baseline load?
upvoted 1 times

Question #48 *Topic 1*

Your organization is developing a plan for migrating to Google Cloud.
What is a best practice when initially configuring your Google Cloud environment?

A. Create a project via Google Cloud Console per department in your company

B. Define your resource hierarchy with an organization node on top

C. Create projects based on team membersأ¢â€™ requests

D. Make every member of your company the project owner

Correct Answer: *B*

Question #49 *Topic 1*

Your organization runs many workloads in different Google Cloud projects, each linked to the same billing account. Each project's workload costs can vary from month to month, but the overall combined cost of all projects is relatively stable. Your organization needs to optimize its cost. What should your organization do?

A. Purchase a commitment per project for each projectι€™s usual minimum

B. Create a billing account per project, and link each project to a different billing account

C. Turn on committed use discount sharing, and create a commitment for the combined usage

D. Move all workloads from all different projects into one single consolidated project

Correct Answer: *B*

Billing contacts -

Reference:

https://cloud.google.com/billing/docs/concepts

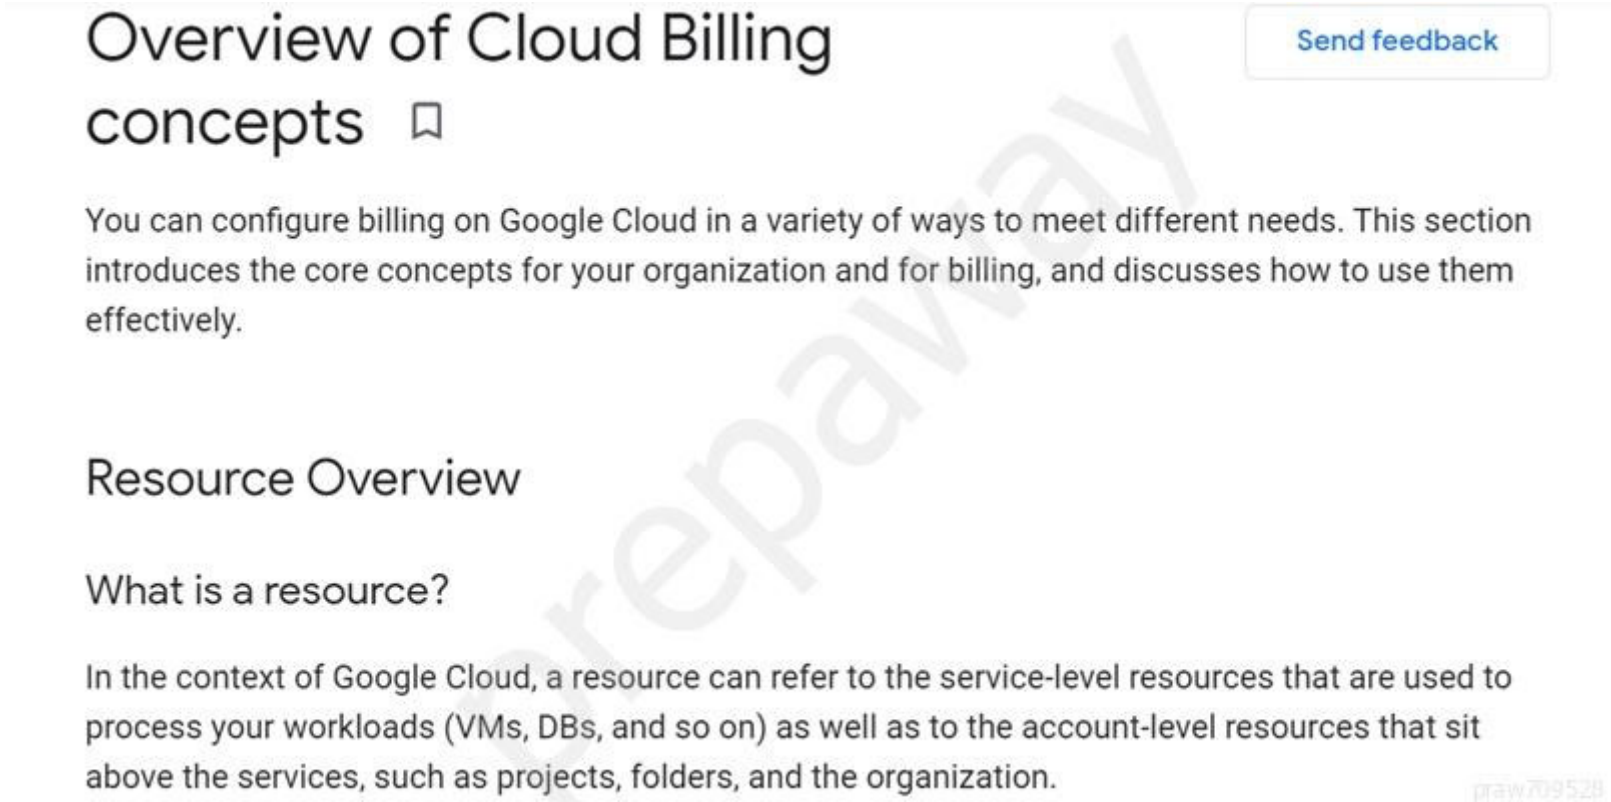
Overview of Cloud Billing concepts

Send feedback

You can configure billing on Google Cloud in a variety of ways to meet different needs. This section introduces the core concepts for your organization and for billing, and discusses how to use them effectively.

Resource Overview

What is a resource?

In the context of Google Cloud, a resource can refer to the service-level resources that are used to process your workloads (VMs, DBs, and so on) as well as to the account-level resources that sit above the services, such as projects, folders, and the organization.

fpreli 1 week, 4 days ago

This is wrong, linking each project to a different billing account won't save customer money. I would say C or A, where C is easier and makes more sense for the specific use case.

upvoted 2 times

ahsangh 1 week, 4 days ago

Selected Answer: C

definitely C for discount sharing across different projects under the same billing account for the same organization

C. Turn on committed use discount sharing, and create a commitment for the combined usage

upvoted 1 times

Question #50 *Topic 1*

How should a multinational organization that is migrating to Google Cloud consider security and privacy regulations to ensure that it is in compliance with global standards?

A. Comply with data security and privacy regulations in each geographical region

B. Comply with regional standards for data security and privacy, because they supersede all international regulations

C. Comply with international standards for data security and privacy, because they supersede all regional regulations

D. Comply with regional data security regulations, because theyλ€™re more complex than privacy standards

Correct Answer: *C*

ahsangh 1 week, 4 days ago
international standards do not always supersede regional regulations

A. Comply with data security and privacy regulations in each geographical region
upvoted 4 times

Question #51 *Topic 1*

Your organization wants to optimize its use of Google Cloudλ€™s discounts on virtual machine-based workloads. You plan to use 200 CPUs constantly for the next 3 years, and you forecast that spikes of up to 300 CPUs will occur approximately 30% of the time. What should you choose?

A. 1-year committed use discount for 200 CPUs

B. 3-year committed use discount for 300 CPUs

C. 3-year committed use discount for 200 CPUs

D. Regular pay-as-you-go pricing

Correct Answer: *D*
Reference:
https://cloud.google.com/compute/all-pricing

fpreli 1 week, 4 days ago
C for me, D won't optimize cost.
upvoted 2 times

ahsangh 1 week, 4 days ago
C. 3-year committed use discount for 200 CPUs
upvoted 1 times

mvdm11111 1 week, 5 days ago
I think here it is C. 3-year committed use discount for 200 CPUs
upvoted 2 times

kiranm2021 1 week, 5 days ago
This doesn't seem right to me . Any feedback or clarification ?
upvoted 1 times

Question #52 *Topic 1*

Your organization needs to minimize how much it pays for data traffic from the Google network to the internet. What should your organization do?

A. Choose the Standard network service tier.

B. Choose the Premium network service tier.

C. Deploy Cloud VPN.

D. Deploy Cloud NAT.

Correct Answer: *C*

Reference:

https://cloud.google.com/blog/products/networking/networking-cost-optimization-best-practices

fpreli 1 week, 4 days ago

I would go with A but I think this question is kind of vague. The link shared in the solution says that evaluating the right Network tier is a possible way to reduce cost.

upvoted 2 times

Question #53 *Topic 1*

Your organization wants to migrate your on-premises environment to Google Cloud. The on-premises environment consists of containers and virtual machine instances. Which Google Cloud products can help to migrate the container images and the virtual machine disks?

A. Compute Engine and Filestore

B. Artifact Registry and Cloud Storage

C. Dataflow and BigQuery

D. Pub/Sub and Cloud Storage

Correct Answer: *A*

Reference:

https://cloud.google.com/compute/docs/import/importing-virtual-disks

If you have virtual disks in your on-premises environment with software and configurations that you need (sometimes referred to as *golden disks* or *golden images*), you can save time by importing those virtual disks into Compute Engine and using the resulting image to create virtual machines. The import tool supports most virtual disk file formats, including VMDK and VHD.

If you exported your disk from Compute Engine, you can create images from the disk.

For information about how to create an automated system for migrating several virtual machines (VMs), see Migrating VMs to Compute Engine.

fpreli 1 week, 4 days ago

I think it's B since Artifact Registry can manage container image migration and Cloud Storage can store virtual disk image.

upvoted 1 times

ahsangh 1 week, 4 days ago

Selected Answer: A

A. Compute Engine

https://cloud.google.com/compute/docs/import

upvoted 1 times

Question #54 *Topic 1*

Your company security team manages access control to production systems using an LDAP directory group.
How is this access control managed in the Google Cloud production project?

A. Assign the proper role to the Service Account in the projectλ€™s IAM Policy

B. Grant each user the roles/iam.serviceAccountUser role on a service account that exists in the Google Group.

C. Assign the proper role to the Google Group in the projectλ€™s IAM Policy.

D. Create the project in a folder with the same name as the LDAP directory group.

Correct Answer: *C*
Reference:
https://cloud.google.com/blog/products/identity-security/achieving-identity-and-access-governance-on-google-cloud

Question #55 *Topic 1*

Your organization wants to be sure that is expenditures on cloud services are in line with the budget. Which two Google Cloud cost management features help your organization gain greater visibility into its cloud resource costs? (Choose two.)

A. Billing dashboards

B. Resource labels

C. Sustained use discounts

D. Financial governance policies

E. Payments profile

Correct Answer: *BD*
Reference:
https://cloud.google.com/cost-management#section-6

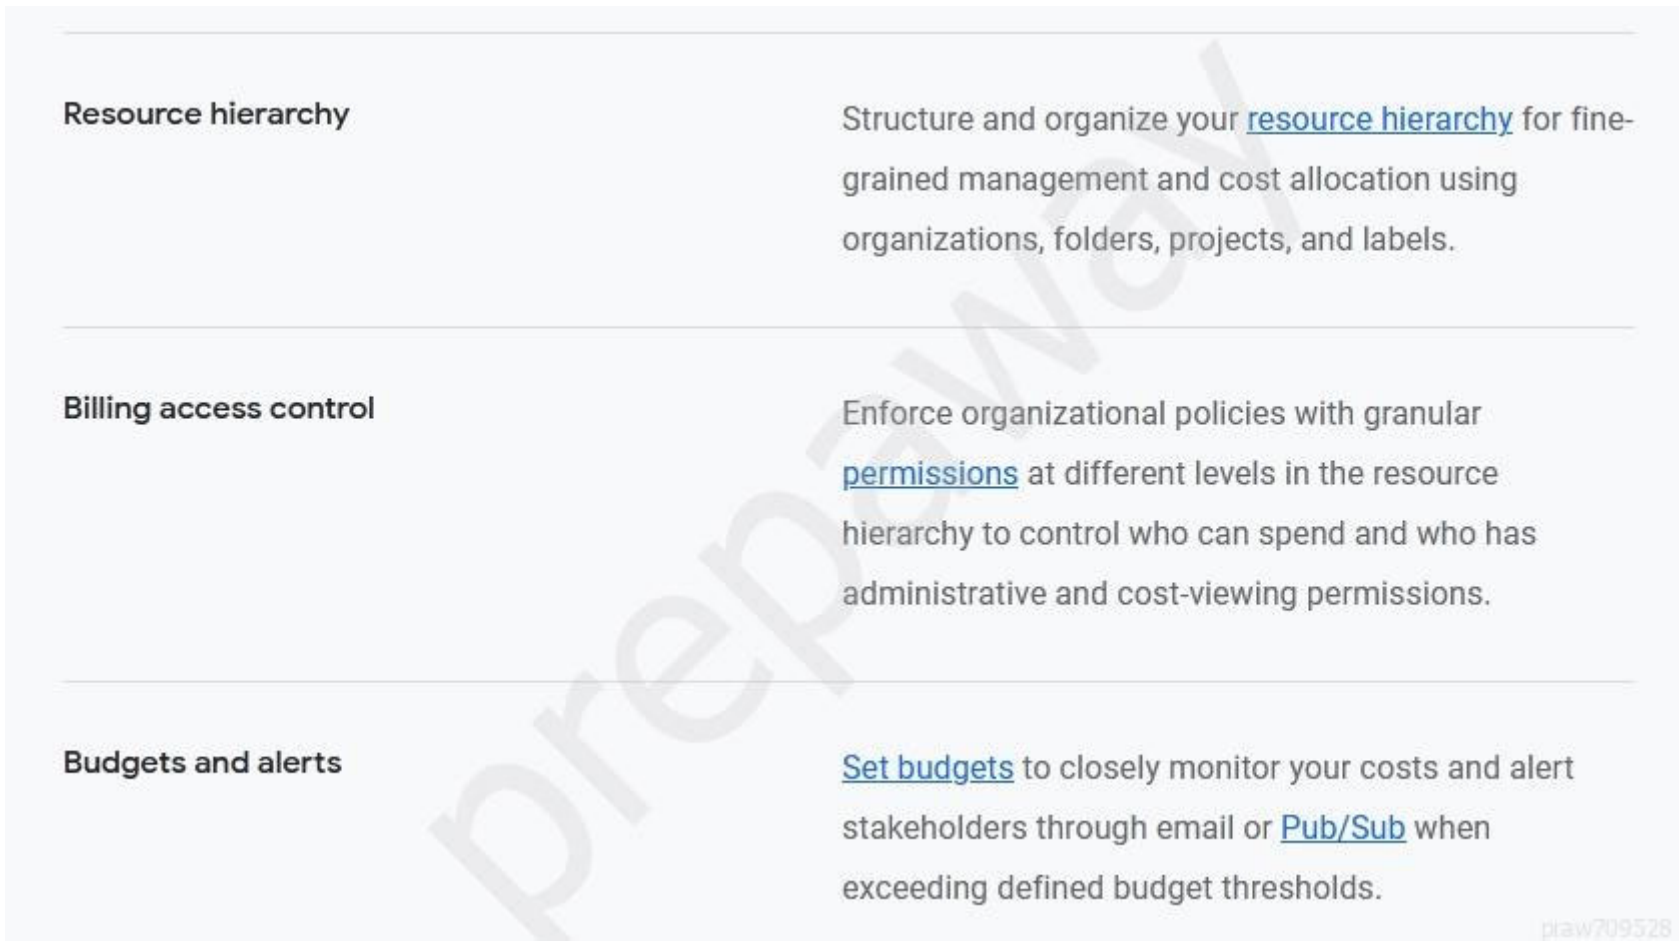

fpreli 1 week, 4 days ago
I think it's AB. Financial governance policies are not related to cost visualization
upvoted 2 times

Question #56 *Topic 1*

Your organization needs to process large amounts of data from an online application that operates continuously. You do not want to be required to provision infrastructure or create server clusters. What should your organization choose?

A. Compute Engine with BigQuery

B. Dataproc

C. Google Kubernetes Engine with Cloud Bigtable

D. Dataflow

Correct Answer: *A*

ahsangh 1 week, 4 days ago
why not D. Dataflow for continuous data
upvoted 1 times

mvdm11111 1 week, 4 days ago
Here it is D
upvoted 4 times

Question #57 *Topic 1*

Your organization needs to ensure that the Google Cloud resources of each of your departments are segregated from one another. Each department has several environments of its own: development, testing, and production. Which strategy should your organization choose?

A. Create a project per department, and create a folder per environment in each project.

B. Create a folder per department, and create a project per environment in each folder.

C. Create a Cloud Identity domain per department, and create a project per environment in each domain.

D. Create a Cloud Identity domain per environment, and create a project per department in each domain.

Correct Answer: *D*
Reference:
https://cloud.google.com/identity/docs/setup

ahsangh 1 week, 4 days ago
Selected Answer: A
A. Create a project per department, and create a folder per environment in each project.
upvoted 1 times

ahsangh 1 week ago
correction: B. Create a folder per department, and create a project per environment in each folder.

https://cloud.google.com/resource-manager/docs/cloud-platform-resource-hierarchy
upvoted 5 times

Question #58 *Topic 1*

Your organization is defining the resource hierarchy for its new application in Google Cloud. You need separate development and production environments. The production environment will be deployed in Compute Engine in two regions. Which structure should your organization choose?

A. Create a single project for all environments. Use labels to segregate resources by environment.

B. Create a single project for all environments. Use tags to segregate resources by environment.

C. Create one project for the development environment and one project for the production environment.

D. Create two projects for the development environment and two projects for the production environment (one for each region).

Correct Answer: *D*

nydo_83 4 days, 8 hours ago
agree with C.
upvoted 1 times

fpreli 5 days, 6 hours ago
C is easier and you won't get less than D. I think creating one project per region wouldn't create benefit.
upvoted 2 times

Question #59 *Topic 1*

Your organization meant to purchase a 3-year Committed Use Discount, but accidentally purchased a 1-year Committed Use Discount instead. What should your organization do?

A. Contact your financial institution.

B. Contact Trust and Safety.

C. Contact Cloud Billing Support.

D. Contact Technical Support.

Correct Answer: *B*

kapomony 4 days, 23 hours ago
Selected Answer: C
is a billing error
upvoted 1 times

fpreli 1 week, 4 days ago
I would go with C, I don't think Security and Compliance has nothing to do with CUD
upvoted 2 times

ahsangh 1 week, 4 days ago
C. Contact Cloud Billing Support.
upvoted 1 times

kiranm2021 1 week, 5 days ago
???? Contact Trust and Safety for billing issues ?
upvoted 1 times

Question #60 *Topic 1*

Your organization needs to allow a production job to have access to a BigQuery dataset. The production job is running on a Compute Engine instance that is part of an instance group.
What should be included in the IAM Policy on the BigQuery dataset?

A. The Compute Engine instance group

B. The project that owns the Compute Engine instance

C. The Compute Engine service account

D. The Compute Engine instance

Correct Answer: *A*
Reference:
https://cloud.google.com/compute/docs/instance-groups

Instance groups

Send feedback

An instance group is a collection of virtual machine (VM) instances that you can manage as a single entity.

Compute Engine offers two kinds of VM instance groups, managed and unmanaged:

- **Managed instance groups** (MIGs) let you operate apps on multiple identical VMs. You can make your workloads scalable and highly available by taking advantage of automated MIG services, including: autoscaling, autohealing, regional (multiple zone) deployment, and automatic updating.
- **Unmanaged instance groups** let you load balance across a fleet of VMs that you manage yourself.

kapomony 4 days, 23 hours ago
answer C --> https://cloud.google.com/bigquery/docs/dataset-access-controls
upvoted 1 times

fpreli 1 week, 4 days ago
C also for me, also because MIG uses the Compute Engine default service account for its own operation
upvoted 3 times

mvdm11111 1 week, 4 days ago
It should be C as this is what the service account is about
upvoted 2 times

Question #61 *Topic 1*

Your team is publishing research results and needs to make large amounts of data available to other researchers within the professional community and the public at minimum cost.
How should you host the data?

A. Use a Cloud Storage bucket and enable λ€Requester Pays.λ€

B. Use a Cloud Storage bucket and provide Signed URLs for the data files.

C. Use a Cloud Storage bucket and set up a Cloud Interconnect connection to allow access to the data.

D. Host the data on-premises, and set up a Cloud Interconnect connection to allow access to the data.

Correct Answer: *D*

jvg637 4 days, 3 hours ago

A. Enabling Requester Pays is useful, for example, if you have a lot of data you want to make available to users, but you don't want to be charged for their access to that data.
https://cloud.google.com/storage/docs/requester-pays
upvoted 1 times

ahsangh 1 week, 4 days ago

Selected Answer: B

B. Signed URLs
upvoted 2 times

Question #62 *Topic 1*

Your company needs to segment Google Cloud resources used by each team from the others. The teamsλ€™ efforts are changing frequently, and you need to reduce operational risk and maintain cost visibility. Which approach does Google recommend?

A. One project per team.

B. One organization per team.

C. One project that contains all of each teamλ€™s resources.

D. One top-level folder per team.

Correct Answer: *A*

Reference:

https://cloud.google.com/security/infrastructure/design

nydo_83 4 days, 8 hours ago

Why not D instead.
A. will tie the team to only have one single project
upvoted 1 times

Question #63 *Topic 1*

How do Migrate for Compute Engine and Migrate for Anthos differ?

A. Unlike Migrate for Anthos, Migrate for Compute Engine assumes that the migration source is VMware vSphere.

B. Migrate for Compute Engine charges for ingress, but Migrate for Anthos does not.

C. Migrate for Compute Engine is closed source, and Migrate for Anthos is open source.

D. Migrate for Anthos migrates to containers, and Migrate for Compute Engine migrates to virtual machines.

Correct Answer: *D*

Reference:

https://cloud.google.com/migrate/anthos

Migrate for Anthos and GKE makes it fast and easy to modernize traditional applications away from virtual machines and into native containers. Our unique automated approach extracts the critical application elements from the VM so you can easily insert those elements into containers in Google Kubernetes Engine or Anthos clusters without the VM layers (like Guest OS) that become unnecessary with containers.

This significantly reduces the cost and labor that would be required for a manual application modernization project. Upon completion, it also empowers your teams to deploy, operate, and maintain existing applications more efficiently and cost effectively, using the same services, policies, and methodologies of a modern platform.

Craft your ideal migration journey

Most digital transformations will be a mix of strategies. For the workloads that will benefit from containers, Migrate for Anthos and GKE delivers a fast, smooth path to modernization. For other workloads that are better suited as a VM, simply move them as is with Migrate for Compute Engine and leverage VPC network integration with GKE. Don't settle for being locked into existing infrastructure or one migration path. With Google, run your workloads how you want, where you want.

Question #64 *Topic 1*

Your large and frequently changing organizationλ€™s user information is stored in an on-premises LDAP database. The database includes user passwords and group and organization membership.
How should your organization provision Google accounts and groups to access Google Cloud resources?

A. Replicate the LDAP infrastructure on Compute Engine

B. Use the Firebase Authentication REST API to create users

C. Use Google Cloud Directory Sync to create users

D. Use the Identity Platform REST API to create users

Correct Answer: *C*
You can run a single instance of Google Cloud Directory Sync to synchronize user accounts and groups to Google Cloud.
Reference:
https://cloud.google.com/architecture/identity/federating-gcp-with-active-directory-introduction

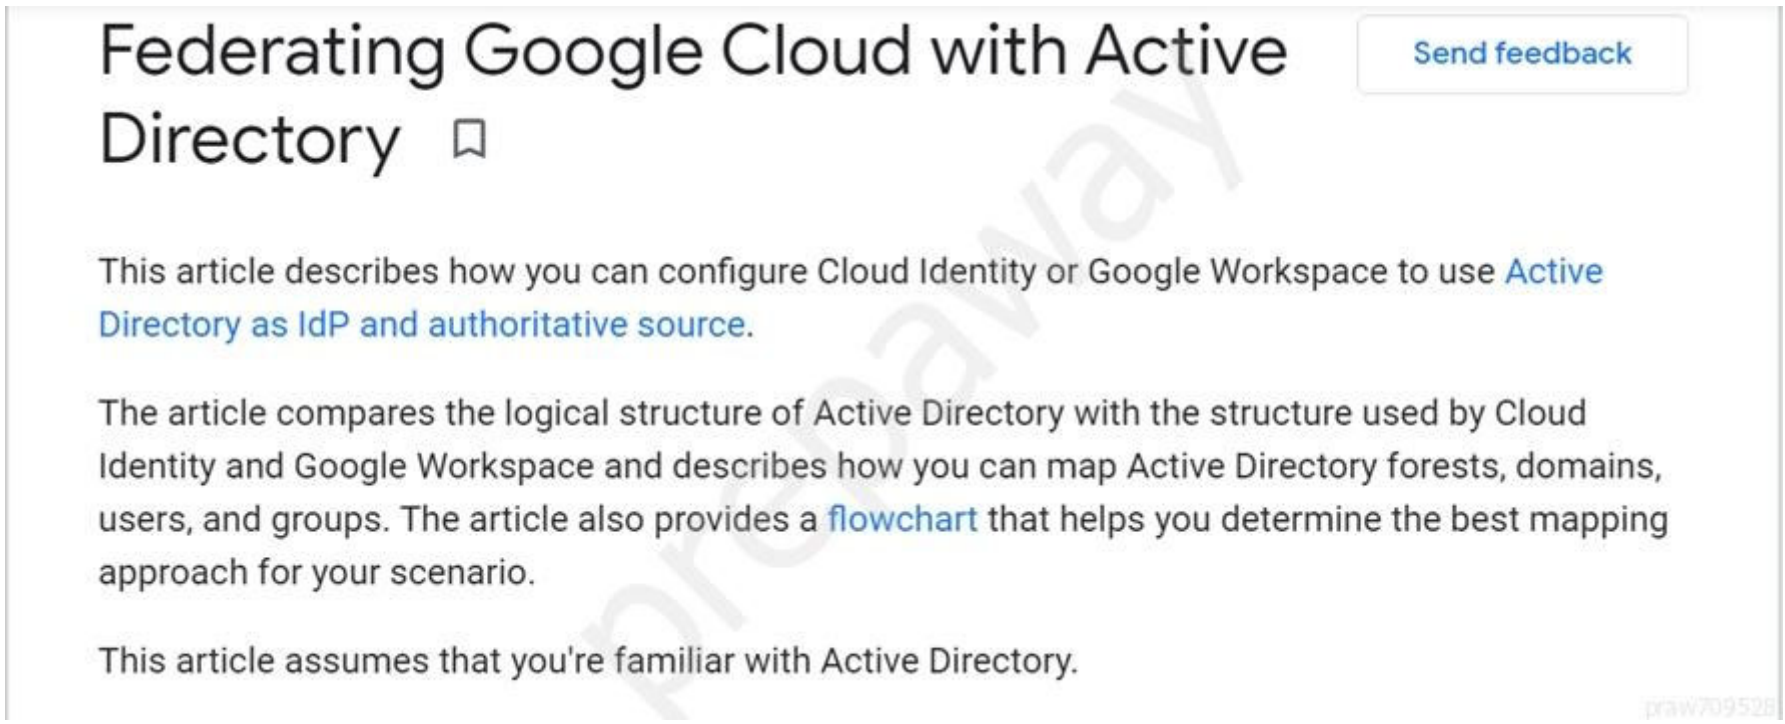

Question #65 *Topic 1*

Your organization recently migrated its compute workloads to Google Cloud. You want these workloads in Google Cloud to privately and securely access your large volume of on-premises data, and you also want to minimize latency.
What should your organization do?

A. Use Storage Transfer Service to securely make your data available to Google Cloud

B. Create a VPC between your on-premises data center and your Google resources

C. Peer your on-premises data center to Googleλ€™s Edge Network

D. Use Transfer Appliance to securely make your data available to Google Cloud

Correct Answer: *B*

Question #66 *Topic 1*

Your organization consists of many teams. Each team has many Google Cloud projects. Your organization wants to simplify the management of identity and access policies for these projects.
How can you group these projects to meet this goal?

A. Group each teamλ€™s projects into a separate domain

B. Assign labels based on the virtual machines that are part of each teamλ€™s projects

C. Use folders to group each teamλ€™s projects

D. Group each teamλ€™s projects into a separate organization node

Correct Answer: *C*